GBLA / LN-tuned CLIP ViT-L/14 — Implemented Inference Pipeline

Tensor shapes on the arrows. Orange = trained parameters (104,450 = deployed θ_s); blue = frozen.

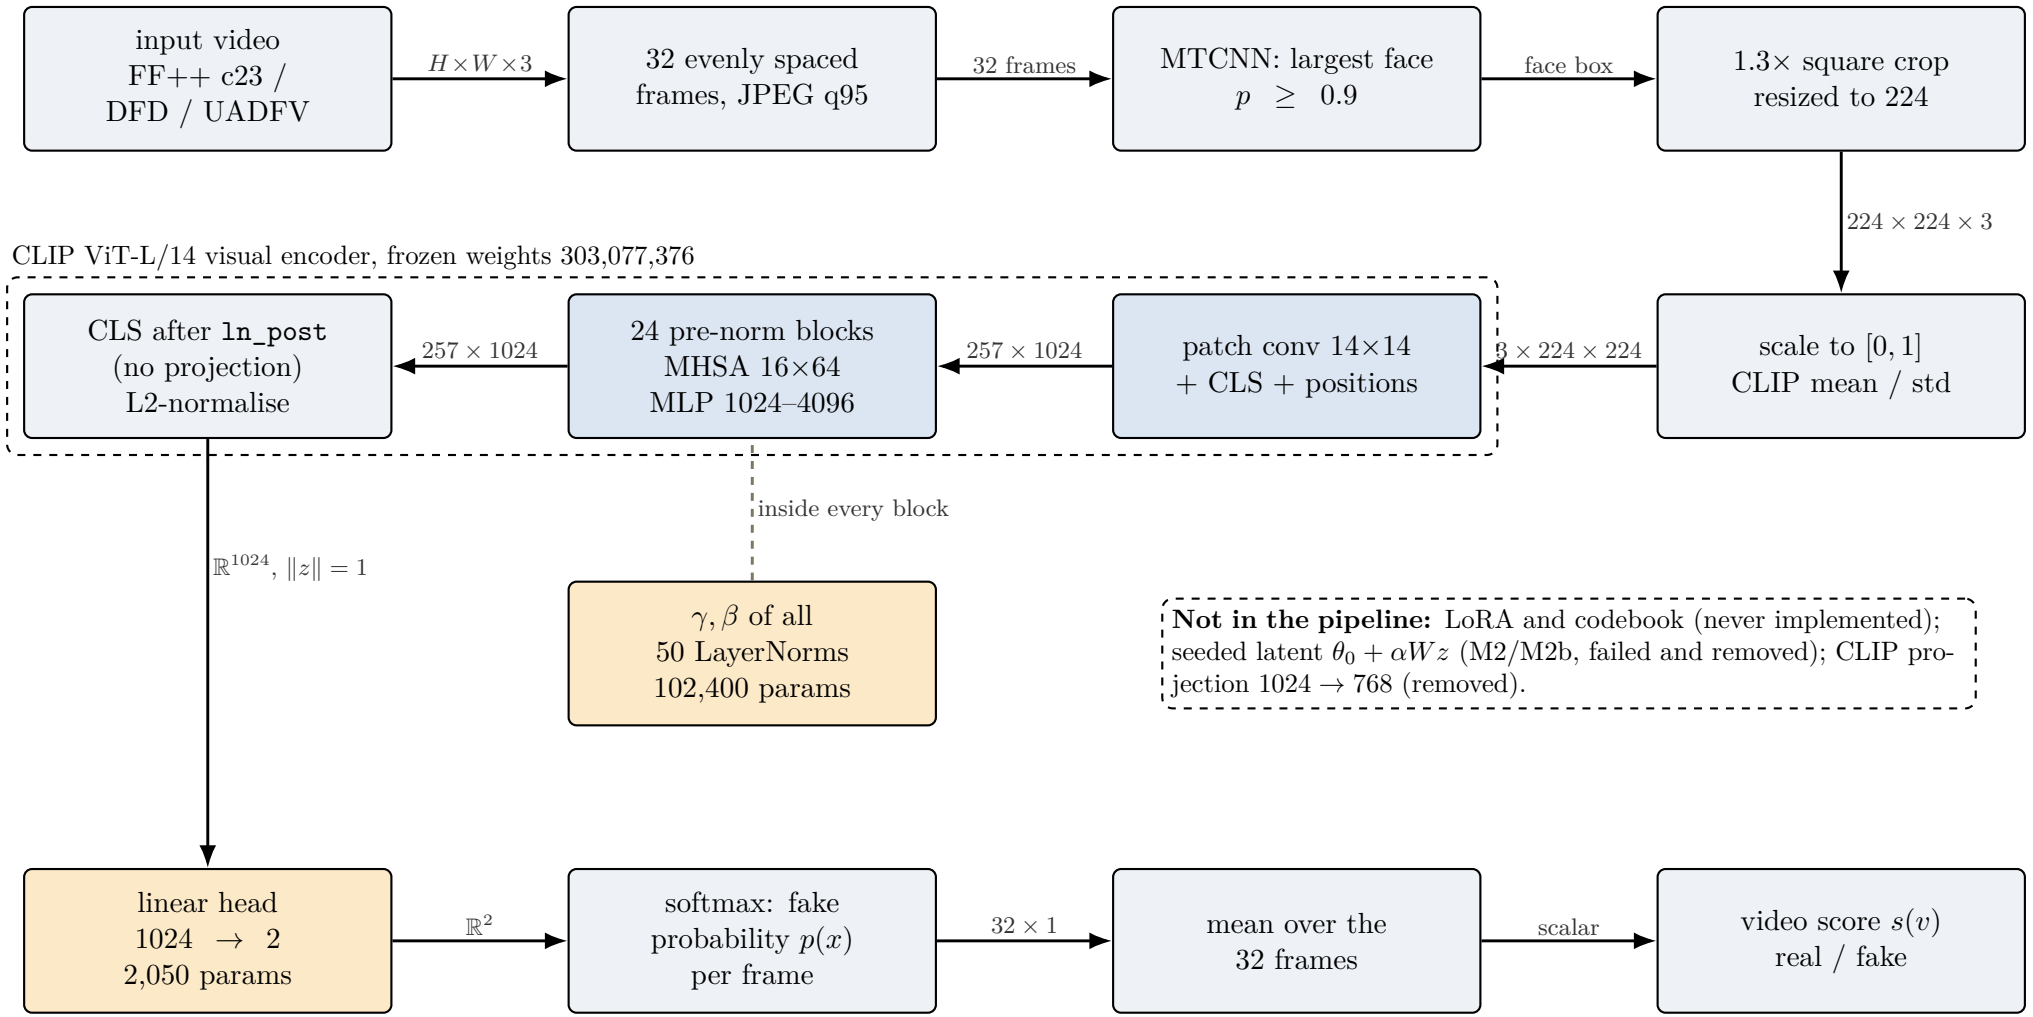

During training, augmentation (flip, affine, blur, colour jitter, JPEG 40–100) is applied between cropping and normalisation. Video score $s(v) = \frac{1}{32} \sum_{t=1}^{32} p(x_t)$.

GBLA Training Workflow (M4 onward)

Shared θ_s plus private offsets δ_g per generator; only θ_s (104,450 numbers) is deployed.

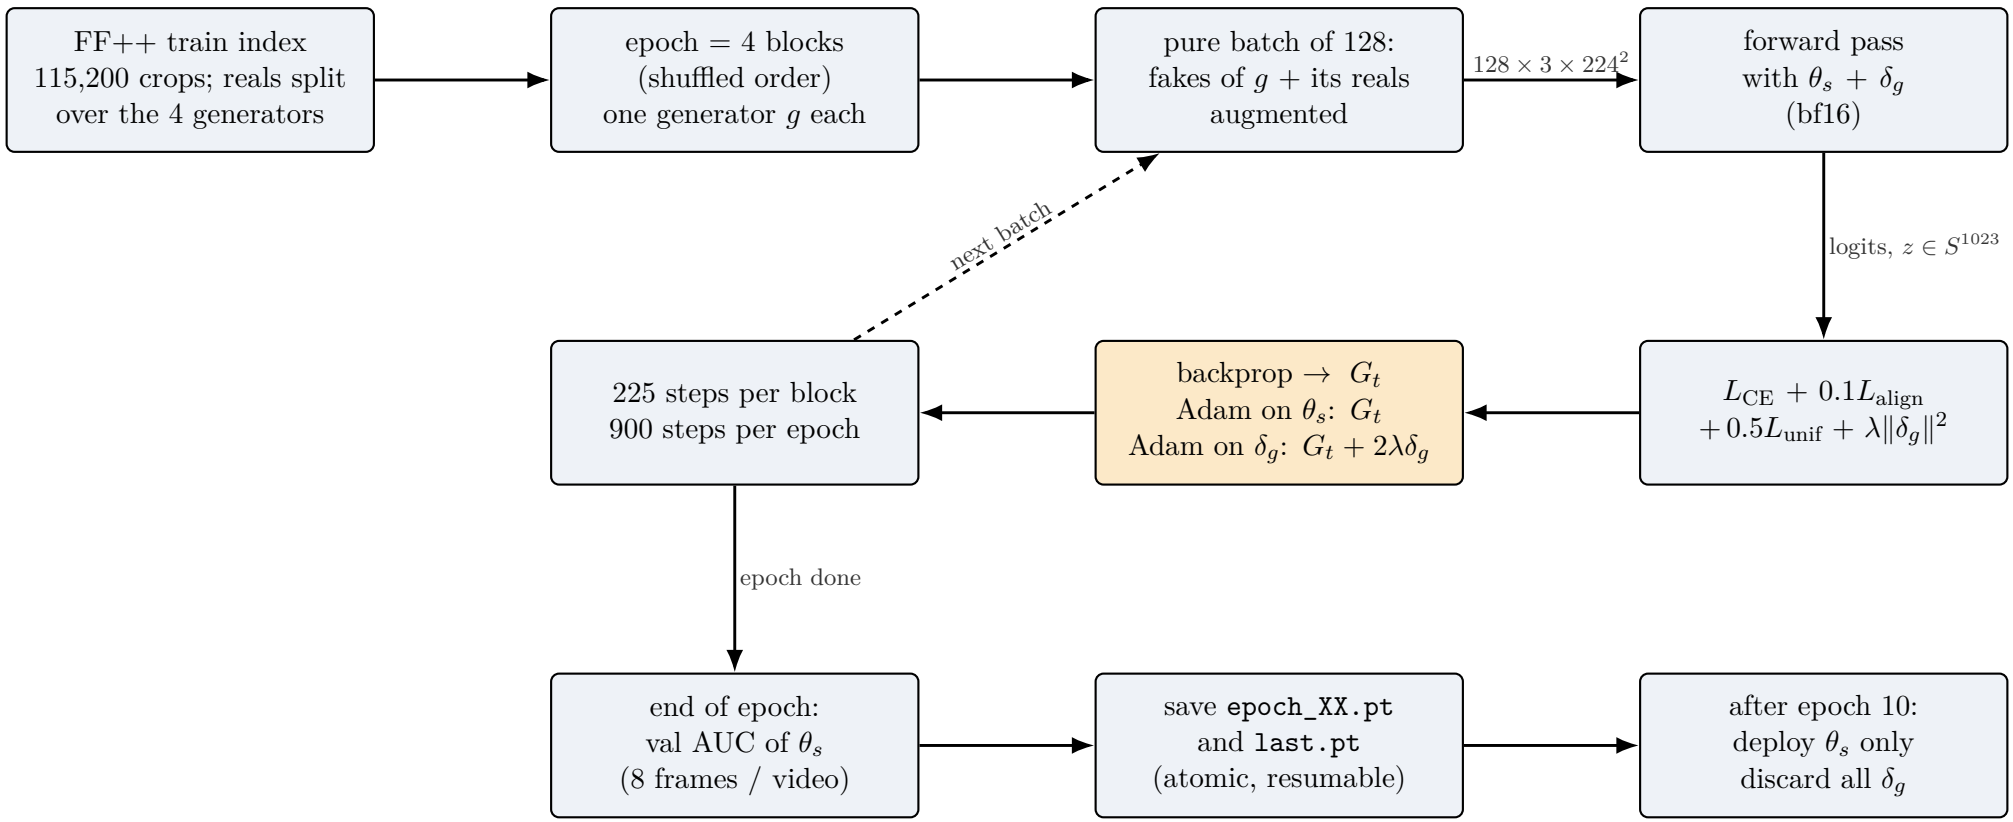

With $\lambda = \infty$ the offsets stay zero, which is pooled training with pure-generator batches; M1 (the LNCLIP-DF baseline) used mixed batches and no offsets. Chosen $\lambda = 0.01$. Optimiser AdamW (no weight decay), learning rate 1 warm-up epoch $10^{-5} \rightarrow 3 \times 10^{-4}$, cosine back to 10^{-5}. No checkpoint is selected by validation or test performance.